

IMPORTING THE LIBRARIES

```
import pandas as pd
from pandas import read_csv
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import LabelEncoder
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import
accuracy_score, classification_report, confusion_matrix, r2_score
from sklearn.tree import DecisionTreeClassifier
from sklearn.naive_bayes import GaussianNB
import warnings
warnings.filterwarnings("ignore")
```

LOADING THE DATASET

```
data=pd.read_csv(r"C:\Users\Admin\Downloads\
1651277648862_healthinsurance.csv")
data.head()
```

	age	sex	weight	bmi	hereditary_diseases	no_of_dependents
smoker \						
0	60.0	male	64	24.3	NoDisease	1
0						
1	49.0	female	75	22.6	NoDisease	1
0						
2	32.0	female	64	17.8	Epilepsy	2
1						
3	61.0	female	53	36.4	NoDisease	1
1						
4	19.0	female	50	20.6	NoDisease	0
0						

	city	bloodpressure	diabetes	regular_ex	job_title	
claim						
0	NewYork		72	0	0	Actor
13112.6						
1	Boston		78	1	1	Engineer
9567.0						
2	Phildelphia		88	1	1	Academician
32734.2						

3	Pittsburg	72	1	0	Chef
48517.6					
4	Buffalo	82	1	0	HomeMakers
1731.7					

```
data.shape
```

```
(15000, 13)
```

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 15000 entries, 0 to 14999
```

```
Data columns (total 13 columns):
```

#	Column	Non-Null Count	Dtype
0	age	14604 non-null	float64
1	sex	15000 non-null	object
2	weight	15000 non-null	int64
3	bmi	14044 non-null	float64
4	hereditary_diseases	15000 non-null	object
5	no_of_dependents	15000 non-null	int64
6	smoker	15000 non-null	int64
7	city	15000 non-null	object
8	bloodpressure	15000 non-null	int64
9	diabetes	15000 non-null	int64
10	regular_ex	15000 non-null	int64
11	job_title	15000 non-null	object
12	claim	15000 non-null	float64

```
dtypes: float64(3), int64(6), object(4)
```

```
memory usage: 1.5+ MB
```

```
data.isnull().sum()
```

age	396
sex	0
weight	0
bmi	956
hereditary_diseases	0
no_of_dependents	0
smoker	0
city	0
bloodpressure	0
diabetes	0
regular_ex	0
job_title	0
claim	0

```
dtype: int64
```

```
data.dropna()
```

smoker	age	sex	weight	bmi	hereditary_diseases	no_of_dependents
0	60.0	1	64	24.3	8	1
0						
1	49.0	0	75	22.6	8	1
0						
2	32.0	0	64	17.8	4	2
1						
3	61.0	0	53	36.4	8	1
1						
4	19.0	0	50	20.6	8	0
0						
...
...						
14995	39.0	1	49	28.3	8	1
1						
14996	39.0	1	74	29.6	8	4
0						
14997	20.0	1	62	33.3	8	0
0						
14998	52.0	1	88	36.7	8	0
0						
14999	52.0	1	57	26.4	8	3
0						
	city	bloodpressure	diabetes	regular_ex	job_title	claim
0	55	72	0	0	2	13112.6
1	5	78	1	1	16	9567.0
2	63	88	1	1	0	32734.2
3	64	72	1	0	10	48517.6
4	8	82	1	0	22	1731.7
...
14995	24	54	1	0	20	21082.2
14996	49	64	1	0	33	7512.3
14997	82	52	1	0	18	1391.5
14998	61	70	1	0	17	9144.6
14999	37	72	1	0	28	25992.8

[13648 rows x 13 columns]

```
data[:,]=np.nan_to_num(data)
data.describe()
```

	age	sex	weight	bmi	\
count	15000.000000	15000.000000	15000.000000	15000.000000	
mean	38.503467	0.489867	64.909600	28.337433	
std	15.213913	0.499914	13.701935	9.474541	
min	0.000000	0.000000	34.000000	0.000000	
25%	26.000000	0.000000	54.000000	25.000000	

50%	40.000000	0.000000	63.000000	28.800000
75%	51.000000	1.000000	76.000000	34.100000
max	64.000000	1.000000	95.000000	53.100000

	hereditary_diseases	no_of_dependents	smoker
city \			
count	15000.000000	15000.000000	15000.000000
mean	7.730533	1.129733	0.198133
std	1.251250	1.228469	0.398606
min	0.000000	0.000000	0.000000
25%	8.000000	0.000000	0.000000
50%	8.000000	1.000000	0.000000
75%	8.000000	2.000000	0.000000
max	9.000000	5.000000	1.000000

	bloodpressure	diabetes	regular_ex	job_title
claim				
count	15000.000000	15000.000000	15000.000000	15000.000000
mean	68.650133	0.777000	0.224133	18.662267
std	19.418515	0.416272	0.417024	10.429298
min	0.000000	0.000000	0.000000	0.000000
25%	64.000000	1.000000	0.000000	10.000000
50%	71.000000	1.000000	0.000000	20.000000
75%	80.000000	1.000000	0.000000	28.000000
max	122.000000	1.000000	1.000000	34.000000

data.keys()

```
Index(['age', 'sex', 'weight', 'bmi', 'hereditary_diseases',
      'no_of_dependents', 'smoker', 'city', 'bloodpressure',
      'diabetes',
      'regular_ex', 'job_title', 'claim'],
      dtype='object')
```

```

le=LabelEncoder()
data['sex']=le.fit_transform(data['sex'])
data['sex'].unique()

data['hereditary_diseases']=le.fit_transform(data['hereditary_diseases'])
data['hereditary_diseases'].unique()

data['city']=le.fit_transform(data['city'])
data['city'].unique()

data['job_title']=le.fit_transform(data['job_title'])
data['job_title'].unique()

array([ 2, 16,  0, 10, 22, 12, 32, 13, 30, 33, 15, 28, 29,  5,  8,  6,
        9,
        26,  1, 19, 34, 18,  4, 23, 20,  7, 31, 14,  3, 11, 24, 17, 25,
       27,
        21], dtype=int64)

```

VISUALIZING THE DATA


```

features =
[['age', 'sex', 'weight', 'bmi', 'smoker', 'bloodpressure', 'diabetes', 'claim']]

plt.subplots(figsize=(20,10))

for i, col in enumerate(features):
    plt.subplot(2,3,i+1)
    sns.distplot(data[col])
plt.show()

```



```
data.corr()
```


	age	sex	weight	bmi	\
age	1.000000	0.033912	0.230385	0.101080	
sex	0.033912	1.000000	0.159249	0.003964	
weight	0.230385	0.159249	1.000000	0.136462	
bmi	0.101080	0.003964	0.136462	1.000000	
hereditary_diseases	0.035083	-0.043909	0.009596	-0.039503	
no_of_dependents	0.020042	0.041440	0.135687	0.017216	
smoker	-0.008754	0.073981	0.015499	0.003260	
city	0.003866	-0.005995	0.013662	-0.002592	
bloodpressure	-0.020699	-0.026718	-0.020835	-0.012007	
diabetes	0.068966	-0.012622	-0.010490	0.085160	
regular_ex	0.013812	0.016332	-0.005578	-0.038384	
job_title	-0.094952	0.001134	-0.070038	-0.010686	
claim	0.294430	0.059592	0.077716	0.098840	
	hereditary_diseases		no_of_dependents		
smoker					\
age	0.035083		0.020042		-0.008754
sex	-0.043909		0.041440		0.073981
weight	0.009596		0.135687		0.015499
bmi	-0.039503		0.017216		0.003260

hereditary_diseases	1.000000	0.018364	-0.390082	
no_of_dependents	0.018364	1.000000	0.008364	
smoker	-0.390082	0.008364	1.000000	
city	-0.004599	-0.002519	0.022218	
bloodpressure	-0.045446	0.024849	0.005709	
diabetes	-0.075312	0.065182	0.058164	
regular_ex	0.016476	-0.010302	-0.036949	
job_title	0.010820	-0.035915	-0.025327	
claim	-0.444337	0.067614	0.773399	
	city	bloodpressure	diabetes	regular_ex
job_title \				
age	0.003866	-0.020699	0.068966	0.013812
0.094952				-
sex	-0.005995	-0.026718	-0.012622	0.016332
0.001134				
weight	0.013662	-0.020835	-0.010490	-0.005578
0.070038				-
bmi	-0.002592	-0.012007	0.085160	-0.038384
0.010686				-
hereditary_diseases	-0.004599	-0.045446	-0.075312	0.016476
0.010820				
no_of_dependents	-0.002519	0.024849	0.065182	-0.010302
0.035915				-
smoker	0.022218	0.005709	0.058164	-0.036949
0.025327				-
city	1.000000	0.016921	-0.002642	-0.002071
0.005045				
bloodpressure	0.016921	1.000000	-0.016498	0.042493
0.017182				
diabetes	-0.002642	-0.016498	1.000000	0.007960
0.019815				
regular_ex	-0.002071	0.042493	0.007960	1.000000
0.005342				
job_title	0.005045	0.017182	0.019815	0.005342
1.000000				
claim	0.013785	0.013742	0.135371	-0.060492
0.026016				-
	claim			
age	0.294430			

```
sex          0.059592
weight       0.077716
bmi          0.098840
hereditary_diseases -0.444337
no_of_dependents 0.067614
smoker       0.773399
city         0.013785
bloodpressure 0.013742
diabetes     0.135371
regular_ex   -0.060492
job_title    -0.026016
claim        1.000000
```

```
plt.figure(figsize=(18,9))
sns.heatmap(data.corr(),annot = True, cmap ="Accent_r")
```





<AxesSubplot:>



```
features = ['age', 'bmi', 'weight', 'bloodpressure']
```

```
plt.subplots(figsize=(20,15))
```

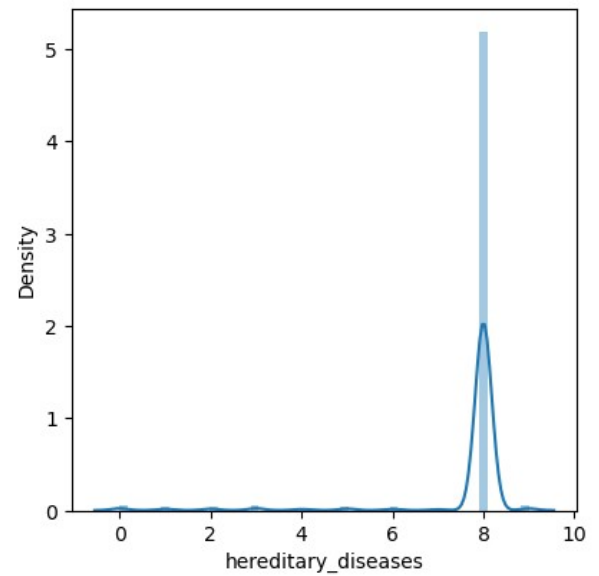



```
for i, col in enumerate(features):
    plt.subplot(2,2,i+1)
    sns.scatterplot(data=data,x=col,y='claim',hue='job_title')
plt.show()
```

```
features = ['age', 'bmi', 'weight', 'hereditary_diseases']
```

```
plt.subplots(figsize=(10,10))
```


```
for i, col in enumerate(features):
    plt.subplot(2,2,i+1)
    sns.distplot(data[col])
plt.show()
```




```
features = ['age', 'bmi', 'weight', 'bmi', 'bloodpressure',]

plt.subplots(figsize=(5,5))


for i, col in enumerate(features):
    plt.subplot(2,3,i+1)
    sns.boxplot(data[col])
plt.show()
```




```
#sns.boxplot(x='age',y='claim',data=data)
sns.lmplot(x='age',y='claim',data=data)
<seaborn.axisgrid.FacetGrid at 0x1da5a086850>
```



```
sns.barplot(x='age',y='claim',data=data)  
<AxesSubplot:xlabel='age', ylabel='claim'>
```



```
sns.jointplot(x='age',y='claim',data=data)
<seaborn.axisgrid.JointGrid at 0x1da59fc4550>
```



```
x=data.iloc[:,0:12].values  
x.shape
```

```
(15000, 12)
```

```
y=data.iloc[:,4].values  
print(y.shape)
```

```
(15000,)
```

Logistic Regression

```
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,random_state=1)
```

```
from sklearn.linear_model import LogisticRegression
reg = LogisticRegression()
reg.fit(x_train,y_train)
```

```
LogisticRegression()
```

```
y_pred=reg.predict(x_test)
```

```
print(classification_report(y_test,y_pred))
print(confusion_matrix(y_test,y_pred))
print("Training Score: ",reg.score(x_train,y_train)*100)
```

	precision	recall	f1-score	support
0	0.69	0.78	0.73	23
1	0.44	0.70	0.54	23
2	0.42	0.19	0.26	26
3	0.40	0.35	0.38	34
4	0.00	0.00	0.00	14
5	0.57	0.59	0.58	27
6	0.00	0.00	0.00	19
7	0.00	0.00	0.00	11
8	0.98	1.00	0.99	2798
9	0.00	0.00	0.00	25
accuracy			0.95	3000
macro avg	0.35	0.36	0.35	3000
weighted avg	0.94	0.95	0.95	3000

```
[[ 18  4  1  0  0  0  0  0  0  0]
 [ 5 16  0  1  0  1  0  0  0  0]
 [ 3  4  5  8  4  1  1  0  0  0]
 [ 0  9  5 12  0  7  1  0  0  0]
 [ 0  3  0  3  0  3  5  0  0  0]
 [ 0  0  1  6  4 16  0  0  0  0]
 [ 0  0  0  0  0  0  0  7 12  0]
 [ 0  0  0  0  0  0  0  0 11  0]
 [ 0  0  0  0  0  0  0  0 2798 0]
 [ 0  0  0  0  0  0  0  0  25  0]]
```

```
Training Score: 96.025
```

```
data = pd.DataFrame({'Actual': y_test, 'Predicted': y_pred})
data
```

	Actual	Predicted
0	8	8
1	8	8
2	8	8
3	8	8
4	9	8
...
2995	8	8
2996	8	8
2997	8	8
2998	8	8
2999	8	8

[3000 rows x 2 columns]

```
print(accuracy_score(y_test,y_pred)*100)
```

95.5

```
from sklearn.model_selection import GridSearchCV
param = {
    'penalty':['l1','l2'],
    'C':[0.001, 0.01, 0.1, 1, 10, 20,100, 1000]
}
lr= LogisticRegression(penalty='l1')
cv=GridSearchCV(reg,param,cv=5,n_jobs=-1)
cv.fit(x_train,y_train)
cv.predict(x_test)
```

array([8, 8, 8, ..., 8, 8, 8], dtype=int64)

```
print("Best CV score", cv.best_score_*100)
```

Best CV score 96.18333333333334

Random Forest

```
ss=StandardScaler()
x_train=ss.fit_transform(x_train)
x_test=ss.transform(x_test)
```

```
clfr=RandomForestClassifier(n_estimators=10,criterion='entropy',random
_state=0)
```

```
clfr.fit(x_train,y_train)
```

```
RandomForestClassifier(criterion='entropy', n_estimators=10,
random_state=0)
```



```
clfr1=RandomForestClassifier(n_estimators=10,criterion='gini',random_state=0)
```

```
clfr1.fit(x_train,y_train)
```

```
RandomForestClassifier(n_estimators=10, random_state=0)
```

```
ypre=clfr.predict(x_test)# entropy ypre calculation
```

```
ypre1=clfr1.predict(x_test)# gini ypre calculation
```

```
print('entropy Accuracy Score:')  
accuracy_score(y_test,ypre)*100
```

```
entropy Accuracy Score:
```

```
99.83333333333333
```

```
print('gini Accuracy Score:')  
accuracy_score(y_test,ypre1)*100
```

```
gini Accuracy Score:
```

```
99.8
```

```
print('entropy - confusion matrix\n-----\n')  
print(confusion_matrix(y_test,ypre))  
print('gini - confusion matrix\n-----\n')  
print(confusion_matrix(y_test,ypre1))
```

```
entropy - confusion matrix
```

```
-----
```

```
[[ 23  0  0  0  0  0  0  0  0  0]  
 [ 1 21  1  0  0  0  0  0  0  0]  
 [ 0  0 26  0  0  0  0  0  0  0]  
 [ 0  0  0 34  0  0  0  0  0  0]  
 [ 0  1  0  0 13  0  0  0  0  0]  
 [ 0  0  0  0  0 27  0  0  0  0]  
 [ 0  0  0  0  0  0 19  0  0  0]  
 [ 0  0  0  0  0  0  0 10  1  0]  
 [ 0  0  0  0  0  0  0  0 2798  0]  
 [ 0  0  0  0  0  0  0  0  1 24]]
```

```
gini - confusion matrix
```

```
-----
```

```
[[ 23  0  0  0  0  0  0  0  0  0]  
 [ 0 23  0  0  0  0  0  0  0  0]]
```

```
[ 0 1 24 1 0 0 0 0 0 0]
[ 0 0 0 34 0 0 0 0 0 0]
[ 0 0 0 2 12 0 0 0 0 0]
[ 0 0 0 0 0 27 0 0 0 0]
[ 0 0 0 0 0 0 19 0 0 0]
[ 0 0 0 0 0 0 1 10 0 0]
[ 0 0 0 0 0 0 0 0 2798 0]
[ 0 0 0 0 0 0 0 0 0 1 24]]
```

```
print('entropy result\n-----')
print(classification_report(y_test,ypre))
print('gini index result\n-----')
print(classification_report(y_test,ypre1))
```

entropy result

	precision	recall	f1-score	support
0	0.96	1.00	0.98	23
1	0.95	0.91	0.93	23
2	0.96	1.00	0.98	26
3	1.00	1.00	1.00	34
4	1.00	0.93	0.96	14
5	1.00	1.00	1.00	27
6	1.00	1.00	1.00	19
7	1.00	0.91	0.95	11
8	1.00	1.00	1.00	2798
9	1.00	0.96	0.98	25
accuracy			1.00	3000
macro avg	0.99	0.97	0.98	3000
weighted avg	1.00	1.00	1.00	3000

gini index result

	precision	recall	f1-score	support
0	1.00	1.00	1.00	23
1	0.96	1.00	0.98	23
2	1.00	0.92	0.96	26
3	0.92	1.00	0.96	34
4	1.00	0.86	0.92	14
5	1.00	1.00	1.00	27
6	0.95	1.00	0.97	19
7	1.00	0.91	0.95	11
8	1.00	1.00	1.00	2798
9	1.00	0.96	0.98	25
accuracy			1.00	3000
macro avg	0.98	0.96	0.97	3000

weighted avg	1.00	1.00	1.00	3000
--------------	------	------	------	------

Decision Tree

```
dtree = DecisionTreeClassifier(max_depth=6, random_state=1)
dtree.fit(x_train,y_train)
DecisionTreeClassifier(max_depth=6, random_state=1)
y_pred=dtree.predict(x_test)
from sklearn.metrics import
classification_report,confusion_matrix,accuracy_score,mean_squared_err
or
print(classification_report(y_test,y_pred))
print(confusion_matrix(y_test,y_pred))
print("Training Score: ",dtree.score(x_train,y_train)*100)
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	23
1	1.00	1.00	1.00	23
2	1.00	1.00	1.00	26
3	1.00	1.00	1.00	34
4	1.00	1.00	1.00	14
5	1.00	1.00	1.00	27
6	1.00	1.00	1.00	19
7	1.00	1.00	1.00	11
8	1.00	1.00	1.00	2798
9	1.00	1.00	1.00	25

	accuracy			
macro avg	1.00	1.00	1.00	3000
weighted avg	1.00	1.00	1.00	3000

[23	0	0	0	0	0	0	0	0	0]
[0	23	0	0	0	0	0	0	0	0]
[0	0	26	0	0	0	0	0	0	0]
[0	0	0	34	0	0	0	0	0	0]
[0	0	0	0	14	0	0	0	0	0]
[0	0	0	0	0	27	0	0	0	0]
[0	0	0	0	0	0	19	0	0	0]
[0	0	0	0	0	0	0	11	0	0]
[0	0	0	0	0	0	0	0	2798	0]
[0	0	0	0	0	0	0	0	0	25]]

Training Score: 100.0

```
print(accuracy_score(y_test,y_pred)*100)
100.0
```

Naive Bayes Theorem

```
Gmodel=GaussianNB() # invocation
Gmodel.fit(x_train,y_train)
train_Gpred=Gmodel.predict(x_train)
test_Gpred=Gmodel.predict(x_test)
train_acc_gau=np.mean(train_Gpred==y_train)
test_acc_gau=np.mean(test_Gpred==y_test)
print('gaussian naive bayes - training
accuracy:',train_acc_gau*100,'%')
print('gaussian nb - testing time accuracy:',test_acc_gau*100,'%')

gaussian naive bayes - training accuracy: 100.0 %
gaussian nb - testing time accuracy: 100.0 %
```